

Build a minimal video platform called “Samtulan” focused on distraction-free and responsible learning.

The platform should include:

1. Core Features:
 - Video player with embedded YouTube videos
 - Category-based browsing (Coding, Science, Fitness, DIY)
 - Search functionality
 - Light/Dark theme toggle
 - Focus mode (hide sidebar and distractions)
 - Watch timer tracking daily viewing time
 - Bookmark / Watch Later feature using localStorage
2. Content Trust System:
 - Each video must have metadata: title, description, category, and optional source
 - If source is provided → mark as “Verified”
 - If no source → mark as “Unverified”
3. Harmful Content Detection:
 - Implement a basic keyword detection system using a predefined list (e.g. hate, violence, fake cure)
 - If harmful keywords are found → mark content as “Sensitive” or “Flagged”
 - Add a report button so users can report content
 - If reports exceed a threshold (e.g. 5) → mark as “Flagged”
4. Trust Score System:
 - Calculate a trust score using:
 - Source presence (+40)
 - Non-flagged content (+30)
 - Likes (+2 per like)
 - Reports (-10 per report)
5. Smart Ranking Algorithm:
 - Sort videos dynamically based on trust score
 - Higher trust score → higher visibility
 - Flagged content → lower ranking
6. UI Indicators:
 - Show badges:
 - ☒ Verified

-  Sensitive
-  Flagged
- Show warning message before playing flagged content

7. Philosophy:

- Do not block content aggressively
- Instead, reduce visibility and add transparency
- Encourage responsible content rather than enforce strict censorship

Use HTML, CSS, and JavaScript for implementation. Keep the UI minimal, clean, and distraction-free.

Build a full-stack minimal video platform called “Samtulan” that looks and feels like a simple YouTube-style interface, but internally uses a hidden trust-based moderation system.

IMPORTANT:

The UI should remain clean and familiar (like YouTube), and users should NOT see the internal moderation logic. All control systems must run invisibly in the background.

CORE FEATURES (Visible to User):

- Video player using YouTube embeds
- Category-based browsing (Coding, Science, Fitness, DIY)
- Search functionality
- Light/Dark theme toggle
- Focus mode (hide sidebar and distractions)
- Watch timer (track daily viewing time)
- Bookmark / Watch Later using localStorage
- Clean, minimal, distraction-free UI

HIDDEN MODERATION SYSTEM (Invisible Layer):

1. Metadata Handling:

Each video should store:

- title
- description
- category
- source (optional)
- likes
- reports
- trustScore
- sensitivity level

2. Trust System Logic:

- If source exists → +40 score
- If content is not flagged → +30 score
- Likes → +2 per like
- Reports → -10 per report

3. Harmful Content Detection:

- Implement keyword detection (e.g. hate, violence, fake cure)
- If detected → mark as “sensitive” or “flagged”
- Add user reporting system
- If reports ≥ 5 → mark as flagged

4. Smart Ranking Algorithm:

- Automatically sort videos by trustScore
- Higher trustScore → higher visibility
- Flagged content → pushed lower
- DO NOT block content aggressively

5. Content Labels (Internal):

- verified
- unverified
- sensitive
- flagged

NOTE: These labels should NOT clutter UI. Only subtle indicators or warnings should be shown when needed.

6. Warning System:

- If content is flagged → show a warning before playback
- Allow user to continue watching

PHILOSOPHY:

- Do not implement strict censorship
- Do not block uploads unnecessarily
- Focus on reducing visibility of harmful content
- Encourage responsible content through ranking and trust

TECH STACK:

- Frontend: HTML, CSS, JavaScript (clean modular structure)
- Backend: Firebase (Authentication + Firestore database)
- Store all video data and trust metrics in database

EXTRA:

- Structure code in a clean, scalable way
- Separate UI logic and moderation logic
- Make it easy to extend with AI APIs later

FINAL GOAL:

To create a platform that appears simple and user-friendly, but internally runs a powerful, invisible content integrity and ranking system.

Build a full-stack minimal video learning platform called “Samtulan” that looks like a simple YouTube-style interface but internally uses an AI-powered trust and moderation system.

The UI must remain clean, minimal, and distraction-free. All intelligent systems must work invisibly in the background.

CORE FEATURES (User-facing):

- Video player with YouTube embeds
- Category browsing (Coding, Science, Fitness, DIY)
- Search functionality
- Light/Dark theme toggle
- Focus mode (distraction-free viewing)
- Watch timer (daily tracking)
- Bookmark / Watch Later (localStorage)
- Clean, minimal UI (no addictive elements)

DATA MODEL (IMPORTANT):

Each video should include:

- title
- description
- category
- source (optional)
- likes
- reports
- trustScore
- sensitivity (safe / sensitive / flagged)
- aiScore (content quality score)
- tags (array)
- difficulty (beginner / intermediate / advanced)
- creatorReputation
- watchTimeAvg

AI/ML FEATURES (CORE SYSTEM):

1. Content Analysis (NLP):
 - Use AI to analyze title + description + transcript
 - Detect harmful, misleading, or low-quality content
 - Use contextual understanding (not just keywords)
2. Speech-to-Text:
 - Convert video audio into text

- Run AI analysis on transcript
3. Toxicity & Harm Detection:
- Integrate a toxicity detection API (like Perspective API)
 - Classify content:
 - safe
 - sensitive
 - flagged
4. AI Quality Scoring:
- Evaluate:
 - clarity
 - educational value
 - structure
 - Generate aiScore (0–100)
5. Trust Score System:
- Calculate trustScore using:
- Source present (+40)
 - Not flagged (+30)
 - Likes (+2 per like)
 - Reports (-10 per report)
 - AI quality score contribution
6. Smart Ranking Algorithm:
- Sort videos using:

trustScore + aiScore + engagement
 - Verified + high-quality content → top
 - Harmful content → lower visibility
7. Learning Recommendation System (Non-addictive):
- Suggest next videos based on:
 - category
 - difficulty progression
 - user watch history
 - Avoid infinite scroll or addictive loops

HIDDEN MODERATION SYSTEM:

- All checks run automatically in backend
- No aggressive blocking

- Reduce visibility of harmful content instead
- Show warning for flagged content before playback

UI INDICATORS (SUBTLE):

- Show minimal badges:
 - Verified
 - Sensitive
- Warning popup for flagged content
- Do not clutter UI with too many labels

TECH STACK:

- Frontend: HTML, CSS, JavaScript
- Backend: Firebase (Auth + Firestore)
- AI APIs:
 - OpenAI (content analysis)
 - Google Speech-to-Text (transcripts)
 - Perspective API (toxicity detection)

ARCHITECTURE:

- Separate UI layer and AI moderation layer
- Store AI results in database
- Make system scalable and modular

PHILOSOPHY:

- Do not restrict users aggressively
- Promote responsible and high-quality content
- Replace engagement-based algorithms with trust-based ranking
- Keep system transparent but not intrusive

FINAL GOAL:

Create a platform that feels simple like YouTube, but internally runs an intelligent AI-powered content integrity system focused on learning, trust, and mental well-being.